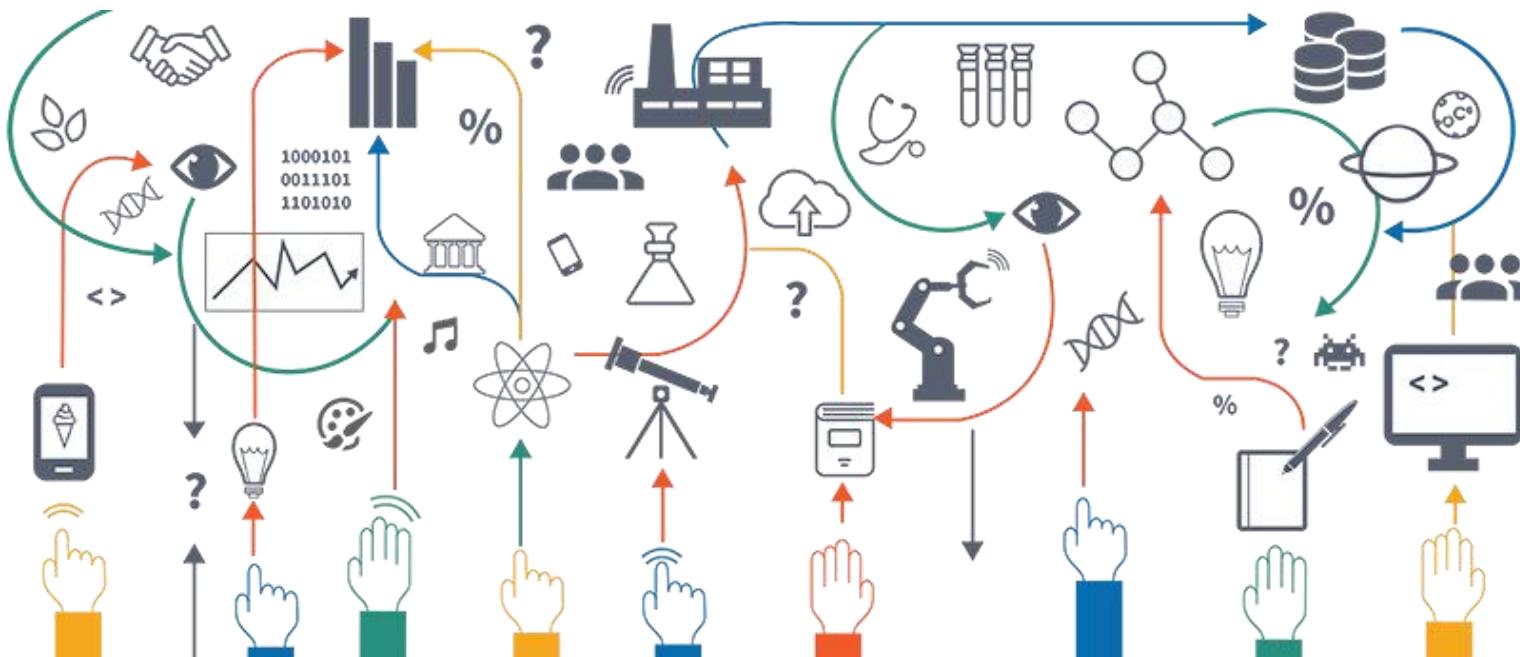


# 6 – Generalization bounds

## Mathematics of Data Science



Lecturer: Dominik Dold



University of Vienna, WiSe 2025  
Master's programme in Data Science

# Motivation

---

In the lecture block on *Function Approximation and Supervised Learning*, we identified the following inequality for the risk:

$$R(\hat{h}) \leq \sup_{h \in \mathcal{H}} |R(h) - \hat{R}(h)| + \inf_{g \in \mathcal{H}} \hat{R}(g)$$

The second term is the **interpolation error**. In the *Deep Neural Networks* block, we introduced the concept of **affine pieces** to bound it!

But how can we bound the first term: **the generalization error?**

# Content

---

- Reminder: Hoeffding's inequality
- PAC learning
- Covering numbers
- Overfitting in the under-and overparametrized regime
- Appendix: VC dimension

# Return of the Hoeffding's inequality

In the lecture block *Foundations of Probability Theory*, we discussed Hoeffding's inequality:

**Hoeffding's inequality:** Assume independent random variables  $x_1, x_2, \dots, x_m$  from the same distribution and finite support  $x \in [a, b]$ . Then:

$$P\left(\left|\frac{1}{m} \sum_i x_i - E(x_1)\right| \geq \epsilon\right) \leq 2 e^{-\frac{2m\epsilon^2}{(b-a)^2}} \quad \forall \epsilon > 0$$

## Observation:

The risk  $R(h)$  is an **expectation value**,

and the empirical risk the corresponding **arithmetic mean**!

If we restrict the loss function to be between, e.g.,  $[0,1]$ , then the risk is an expectation value of a random variable with **finite support**!



**Hoeffding's inequality applies! :)**

# Learning bound: finite hypothesis set

Equipped with this knowledge, let's assume we have a finite hypothesis set  $\mathcal{H}$ . We start with the probability that the supremum is larger than some certain error  $\epsilon > 0$ :

$$p\left(\sup_{h \in \mathcal{H}} |R(h) - \hat{R}(h)| \geq \epsilon\right) \leq \sum_{h \in \mathcal{H}} p(|R(h) - \hat{R}(h)| \geq \epsilon) \leq \sum_{h \in \mathcal{H}} 2e^{-2m\epsilon^2} \leq |\mathcal{H}|2e^{-2m\epsilon^2} = \delta$$

Union bound:  
 $p(\text{largest one } \geq \epsilon) \leq p(\text{at least one } \geq \epsilon)$

Hoeffding

# training samples

With this, we get:

Let  $\mathcal{H}$  be a finite hypothesis set. Then for every  $\delta > 0$ , the following inequality holds with probability at least  $1 - \delta$  for all  $h \in \mathcal{H}$ :

$$|R(h) - \hat{R}(h)| \leq \sqrt{\frac{\ln(|\mathcal{H}|) + \ln\left(\frac{2}{\delta}\right)}{2m}} = \epsilon$$

# PAC Learning

Learning algorithms that have such a learning bound are also known as **PAC learnable**:

**Probably Approximately Correct**

$$1 - \delta \qquad \epsilon$$



We can squeeze out some more by solving  $\epsilon$  for the number of training samples  $m$ :

Let  $\mathcal{H}$  be a finite hypothesis set. Then for every  $\delta > 0$  and  $\epsilon > 0$ , we have for all  $h \in \mathcal{H}$

$$p(|R(h) - \hat{R}(h)| \leq \epsilon) \geq 1 - \delta$$

if we have at least  $m \geq \frac{\ln(|\mathcal{H}|) + \ln(\frac{2}{\delta})}{2\epsilon}$  training samples.

# Problem: what happens in the limit of infinite hypotheses?

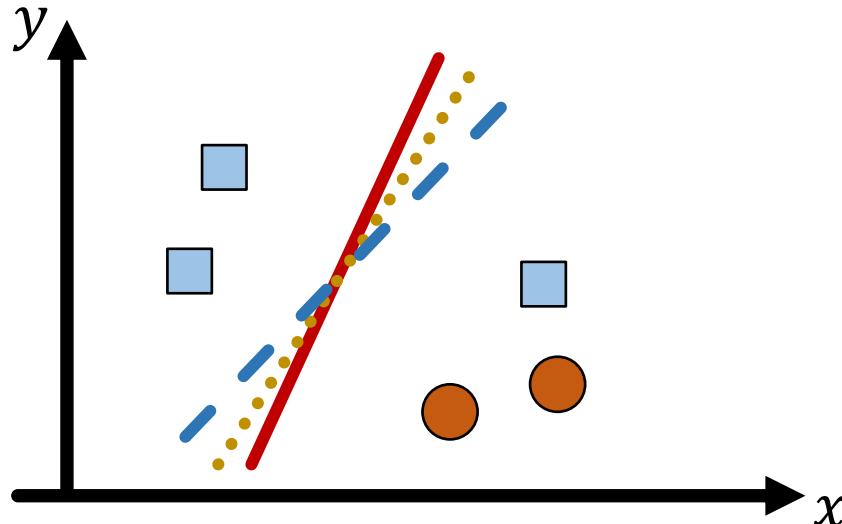
$$\sup_{h \in \mathcal{H}} |R(h) - \hat{R}(h)| \leq \sqrt{\frac{\ln(|\mathcal{H}|) + \ln\left(\frac{2}{\delta}\right)}{2m}}$$

The bound goes to infinity!

**But why?** Obviously linear regression, neural networks, etc. all generalize to some degree!

**Reason:** we assumed independence of the hypotheses! But what if two hypotheses have very similar outputs for all possible inputs? We shouldn't count them separately!

**Illustration:** classification



All lines shown here produce similar outputs!

Thus, instead of counting them individually, we should count them as “one class” of “equivalent” functions!

# Covering number

---

Let's formalize this!

Assume we have our hypothesis set  $\mathcal{H}$ . We will now decompose the set into “equivalence” classes. Per class  $C_k \subset \mathcal{H}$ , we have one representative function  $h_k \in \mathcal{H}$  such that for all  $h \in C_k$ ,  $|h - h_k|_\infty < \kappa$  for  $\kappa > 0$ .

The covering number  $d_C(\kappa)$  is the minimum number of classes  $k$  required to **fully cover**  $\mathcal{H}$ , i.e., such that “ $\bigcup_{i=1}^k C_k = \mathcal{H}$ ” up to error  $\kappa$ .

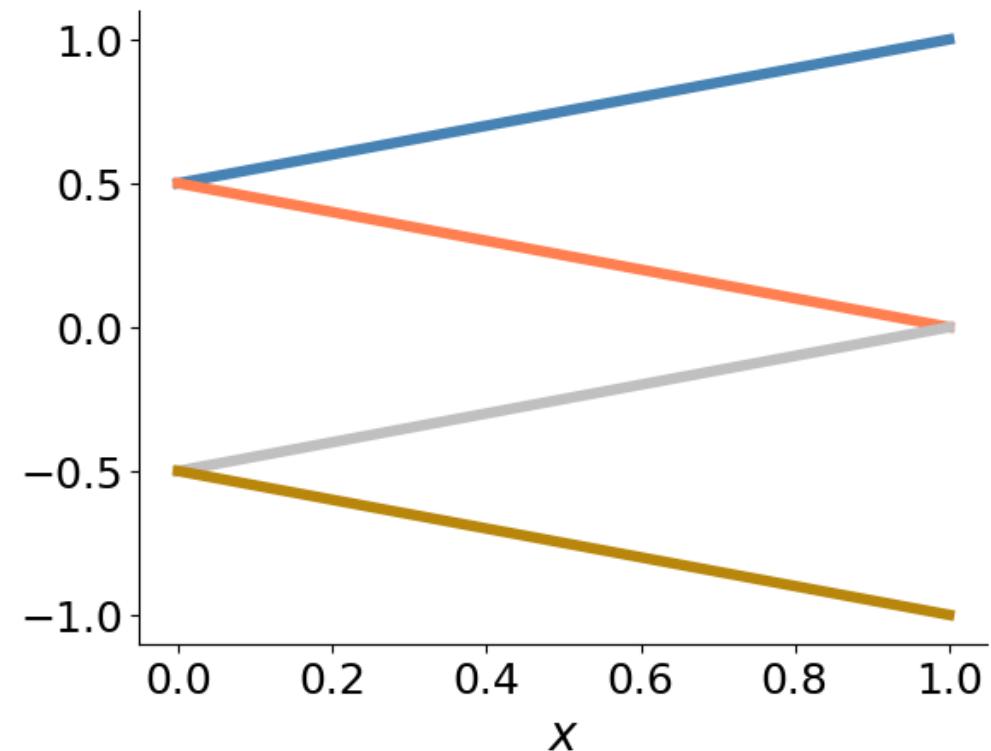
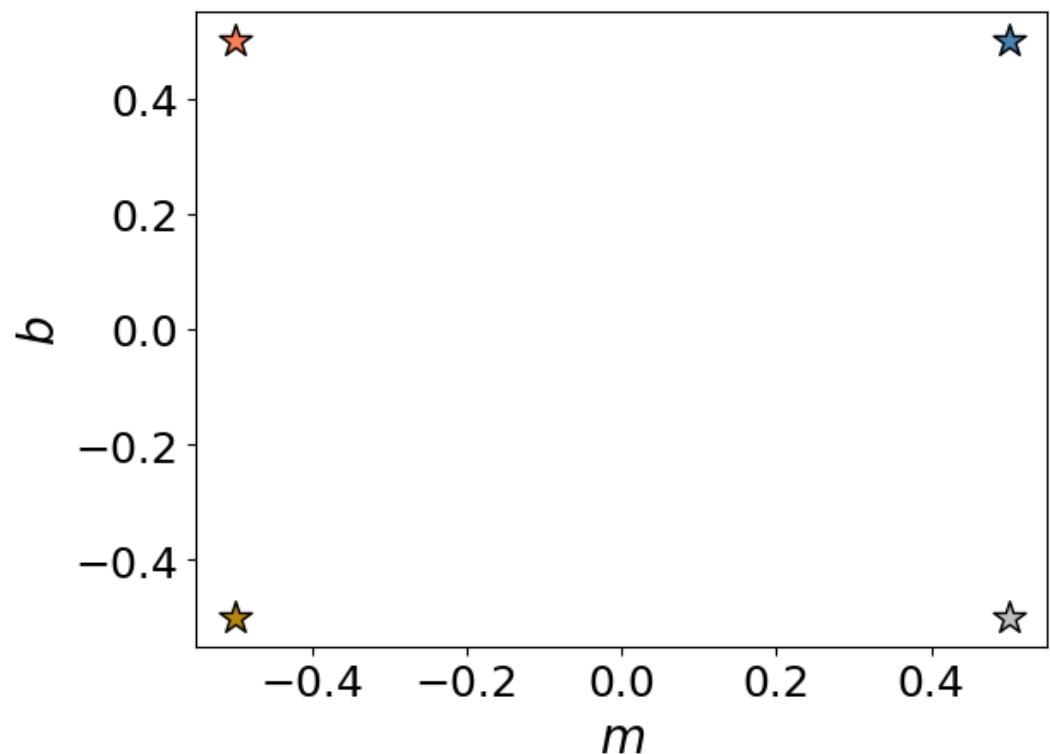
# Example

---

**Linear regression:**  $f_{m,b}(x) = mx + b$  with  $m \in [-0.5, 0.5]$ ,  $b \in [-0.5, 0.5]$ ,  $x \in [0,1]$ .

Assume  $\kappa = 0.5$ . Then the following four functions cover all realizations of  $f$ :

$$f_{0.5,0.5}, f_{-0.5,0.5}, f_{0.5,-0.5}, f_{-0.5,-0.5}$$



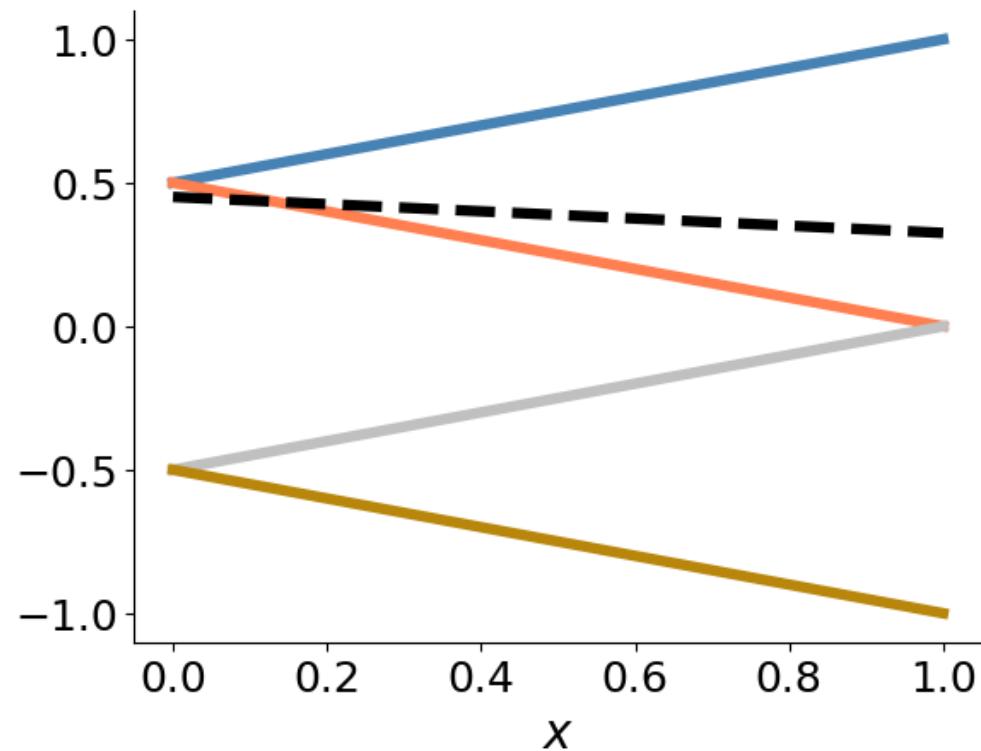
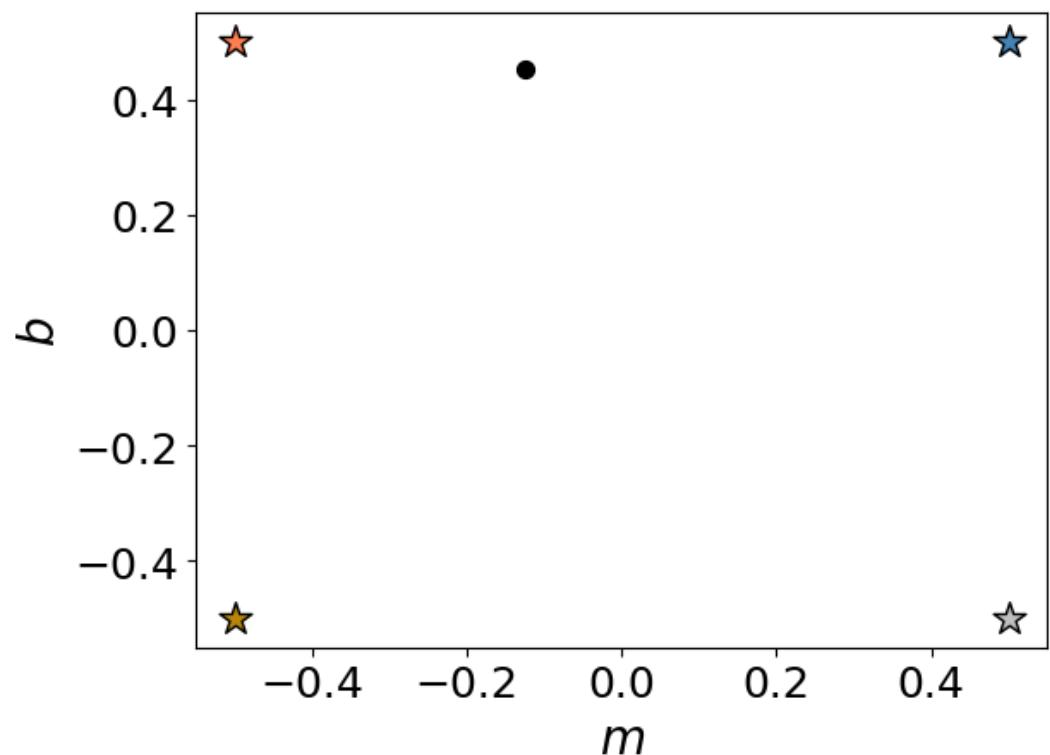
# Example

---

**Linear regression:**  $f_{m,b}(x) = mx + b$  with  $m \in [-0.5, 0.5]$ ,  $b \in [-0.5, 0.5]$ ,  $x \in [0,1]$ .

Assume  $\kappa = 0.5$ . Then the following four functions cover all realizations of  $f$ :

$$f_{0.5,0.5}, f_{-0.5,0.5}, f_{0.5,-0.5}, f_{-0.5,-0.5}$$



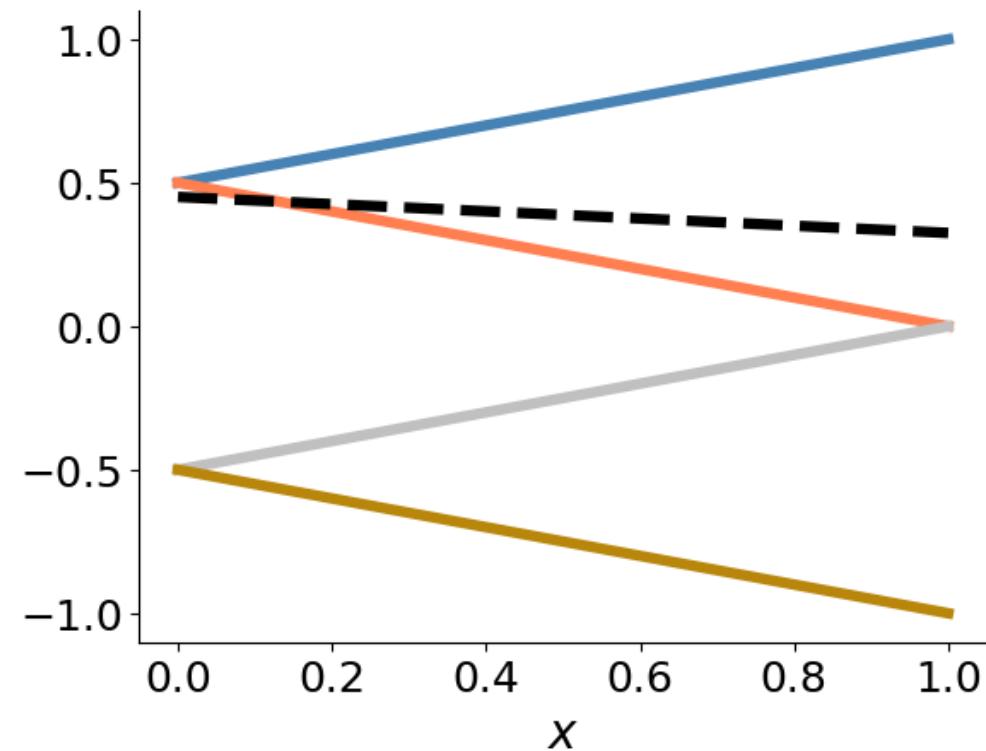
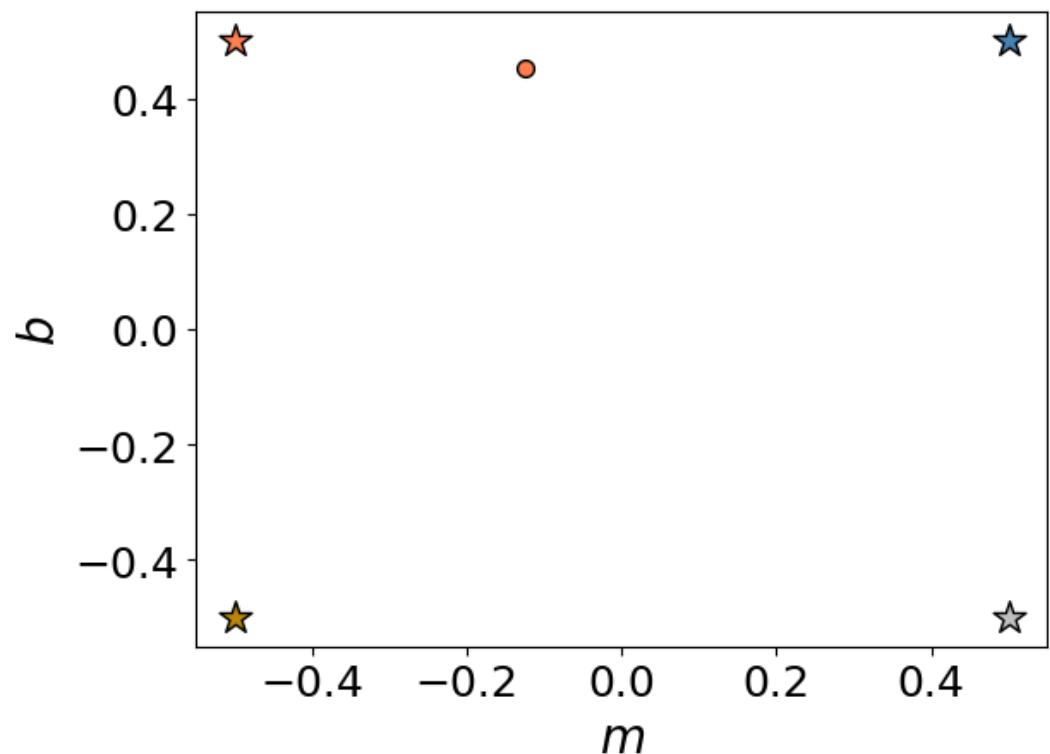
# Example

---

**Linear regression:**  $f_{m,b}(x) = mx + b$  with  $m \in [-0.5, 0.5]$ ,  $b \in [-0.5, 0.5]$ ,  $x \in [0,1]$ .

Assume  $\kappa = 0.5$ . Then the following four functions cover all realizations of  $f$ :

$$f_{0.5,0.5}, f_{-0.5,0.5}, f_{0.5,-0.5}, f_{-0.5,-0.5}$$



# Example

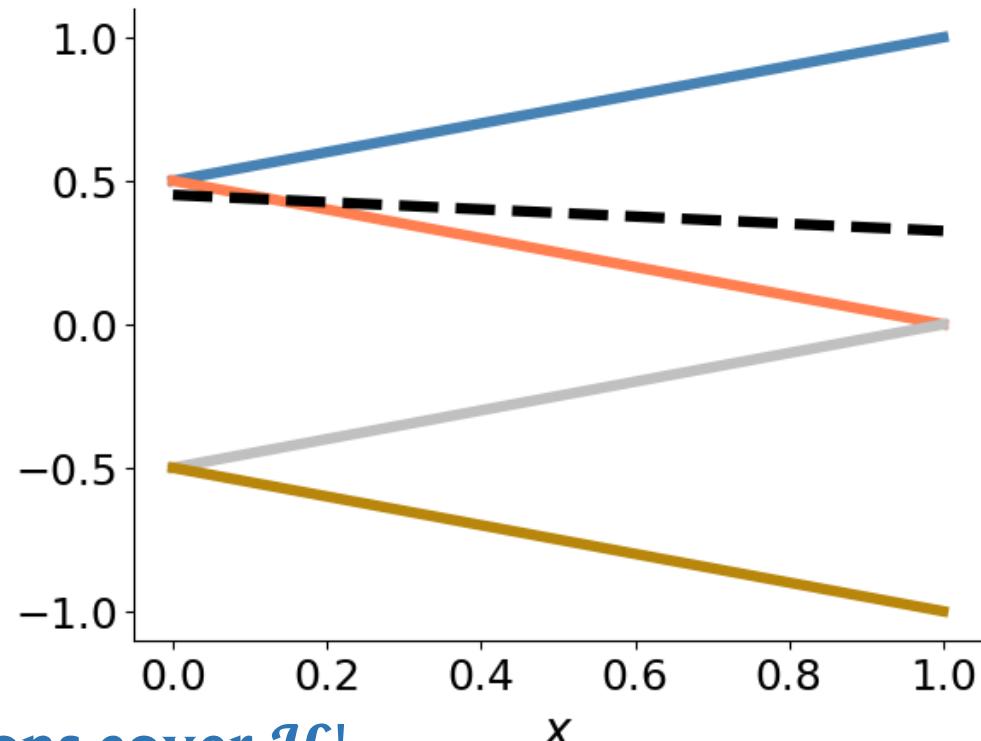
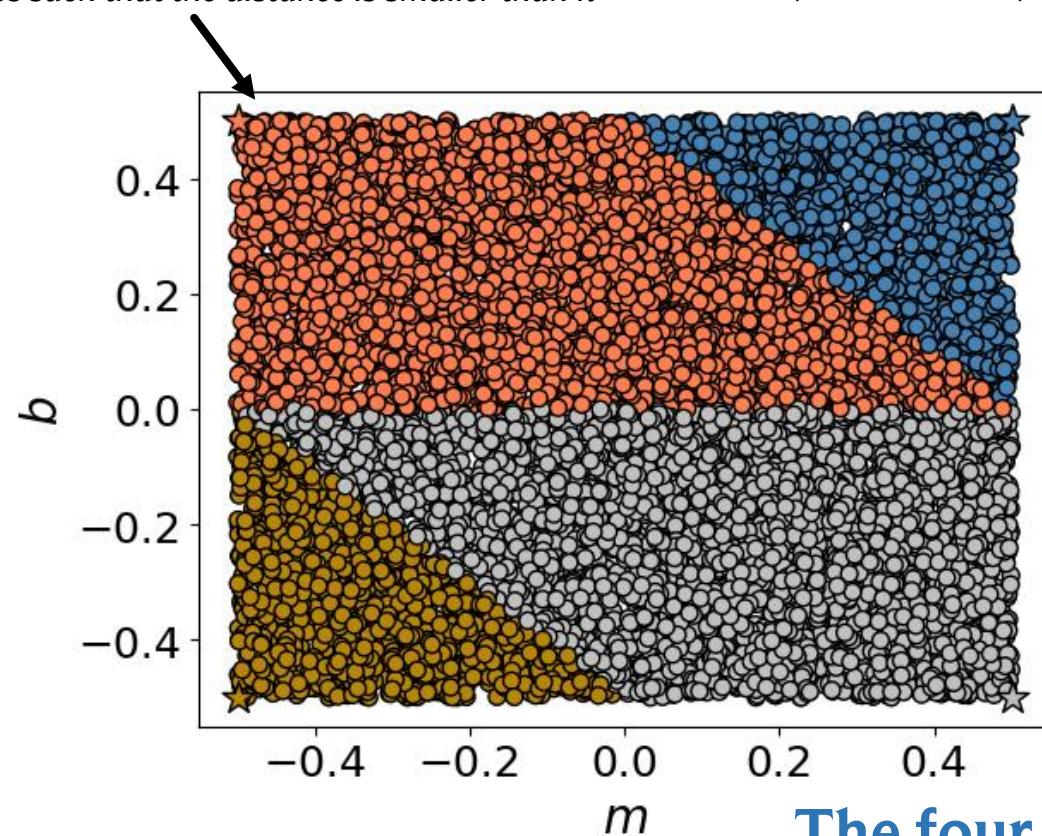
---

**Linear regression:**  $f_{m,b}(x) = mx + b$  with  $m \in [-0.5, 0.5]$ ,  $b \in [-0.5, 0.5]$ ,  $x \in [0,1]$ .

Assume  $\kappa = 0.5$ . Then the following four functions cover all realizations of  $f$ :

*assign random functions to one of the four  
classes such that the distance is smaller than  $\kappa$*

$$f_{0.5,0.5}, f_{-0.5,0.5}, f_{0.5,-0.5}, f_{-0.5,-0.5}$$



The four functions cover  $\mathcal{H}!$

# Learning bound: covering numbers

Previously, we had:

Let  $\mathcal{H}$  be a finite hypothesis set. Then for every  $\delta > 0$ , the following inequality holds with probability at least  $1 - \delta$  for all  $h \in \mathcal{H}$ :

$$|R(h) - \hat{R}(h)| \leq \sqrt{\frac{\ln(|\mathcal{H}|) + \ln(2/\delta)}{2m}}$$

Using covering numbers, we get (*using a rather similar derivation*):

Let  $\mathcal{H}$  be a hypothesis set with covering number  $\mathcal{G}$ . Then for every  $\delta > 0$ , and for  $C_L$ -Lipschitz loss function with output range  $[-C, C]$ , we have with probability at least  $1 - \delta$  for all  $h \in \mathcal{H}$  and  $m$  training data samples:

$$|R(h) - \hat{R}(h)| \leq 4C \cdot C_L \cdot \sqrt{\frac{\ln \mathcal{G} + \ln(2/\delta)}{m}} + \frac{2C_L}{m^\alpha} \text{ size of the "balls" } \kappa \text{ we use for covering } \mathcal{H}$$

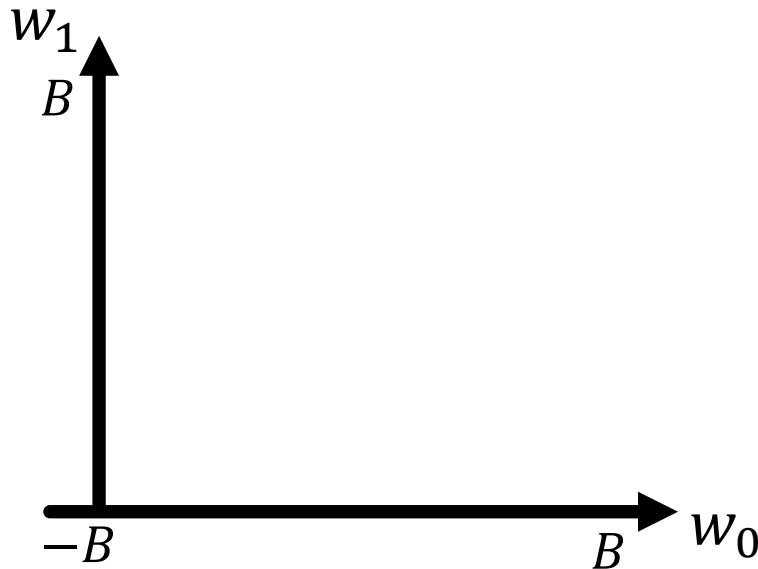
# Covering numbers from Lipschitz constants

---

If our neural network is  **$C$ -Lipschitz in the parameters** (*which it is for most activation functions used in practice*), then there is a clever way of **bounding the covering number!**

Taking the reverse route: cover the parameter space!

→ Find  $w_i$  such that for all  $w \in [-B, B]^n$ , there is at least one  $\alpha_i$  with  $|w - \alpha_i|_\infty \leq \frac{\kappa}{C}$ .

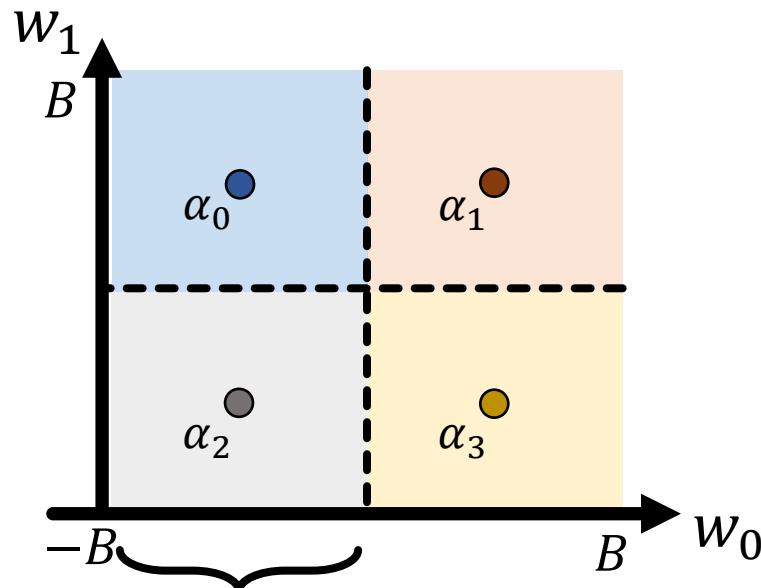


# Covering numbers from Lipschitz constants

If our neural network is  **$C$ -Lipschitz in the parameters** (*which it is for most activation functions used in practice*), then there is a clever way of **bounding the covering number!**

Taking the reverse route: cover the parameter space!

→ Find  $w_i$  such that for all  $w \in [-B, B]^n$ , there is at least one  $\alpha_i$  with  $|w - \alpha_i|_\infty \leq \frac{\kappa}{C}$ .



**Size:**  $\kappa/C$

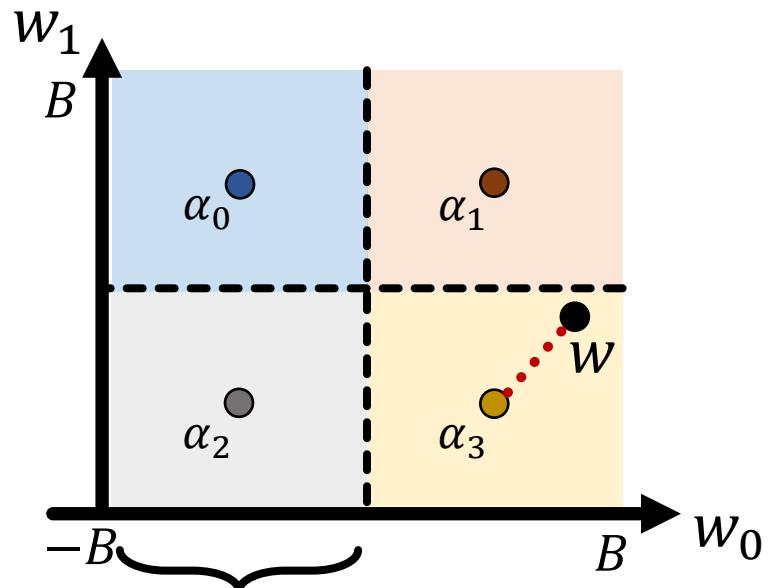
**Thus:**  $\mathcal{G} \leq (B \cdot C/\kappa)^n = \# \text{ of boxes here}$

# Covering numbers from Lipschitz constants

If our neural network is  **$C$ -Lipschitz in the parameters** (which it is for most activation functions used in practice), then there is a clever way of **bounding the covering number!**

Taking the reverse route: cover the parameter space!

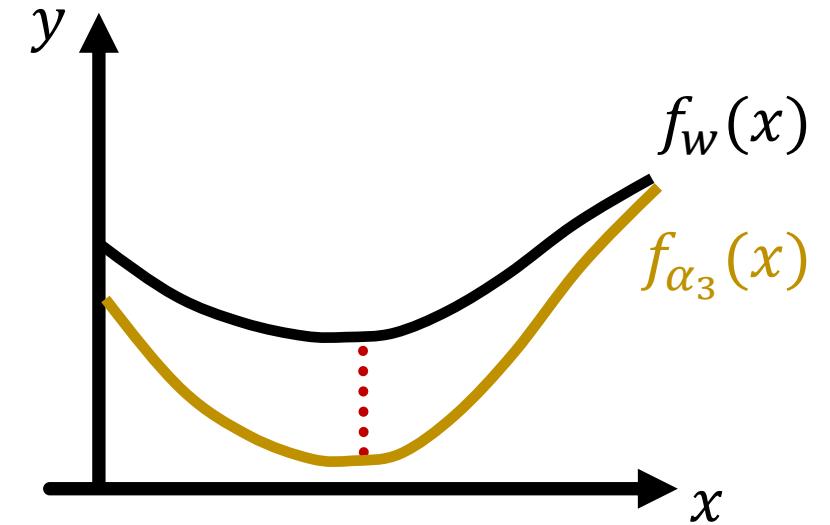
→ Find  $w_i$  such that for all  $w \in [-B, B]^n$ , there is at least one  $\alpha_i$  with  $|w - \alpha_i|_\infty \leq \frac{\kappa}{C}$ .



Size:  $\kappa/C$

Thus:  $\mathcal{G} \leq (B \cdot C/\kappa)^n = \# \text{ of boxes here}$

Look at any  $w$ !



Via Lipschitz property, we have  
 $|f_w(x) - f_{\alpha_i}(x)|_\infty \leq C|w - \alpha_i|_\infty = \kappa$ ,  
→ We also covered  $\mathcal{H}$ !

# Finale: Lipschitz constant of ReLU neural networks

---

A ReLU neural network with  $n$  parameters, max. width  $d$ ,  $L$  layers, an activation function with Lipschitz constant  $C_\phi$ , and parameters constrained to  $[-B, B]$ , has Lipschitz constant:

1. If  $B \geq 1$ :  $C = n \cdot (2 C_\phi B d)^L$
2. If  $B > 0$ :  $C = C_\phi^L \cdot (B d)^{L+1}$

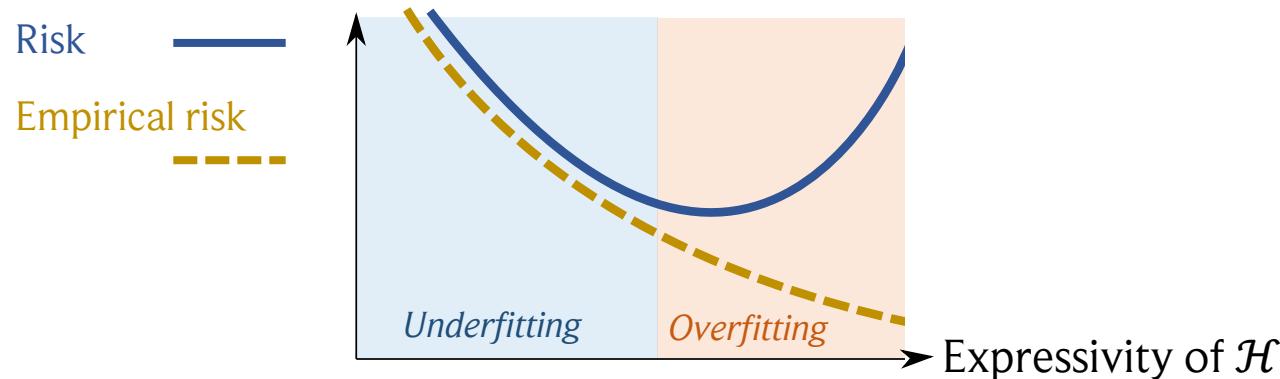
From which we get the covering numbers, e.g.,:

1. If  $B \geq 1$ :  $\mathcal{G} \leq (n/\kappa)^n \cdot (2 C_\phi B d)^{n \cdot L}$
2. If  $B > 0$ :  $\mathcal{G} \leq (C_\phi^L/\kappa)^n \cdot (B d)^{nL+n}$

These can be simply inserted into our generalization bounds!

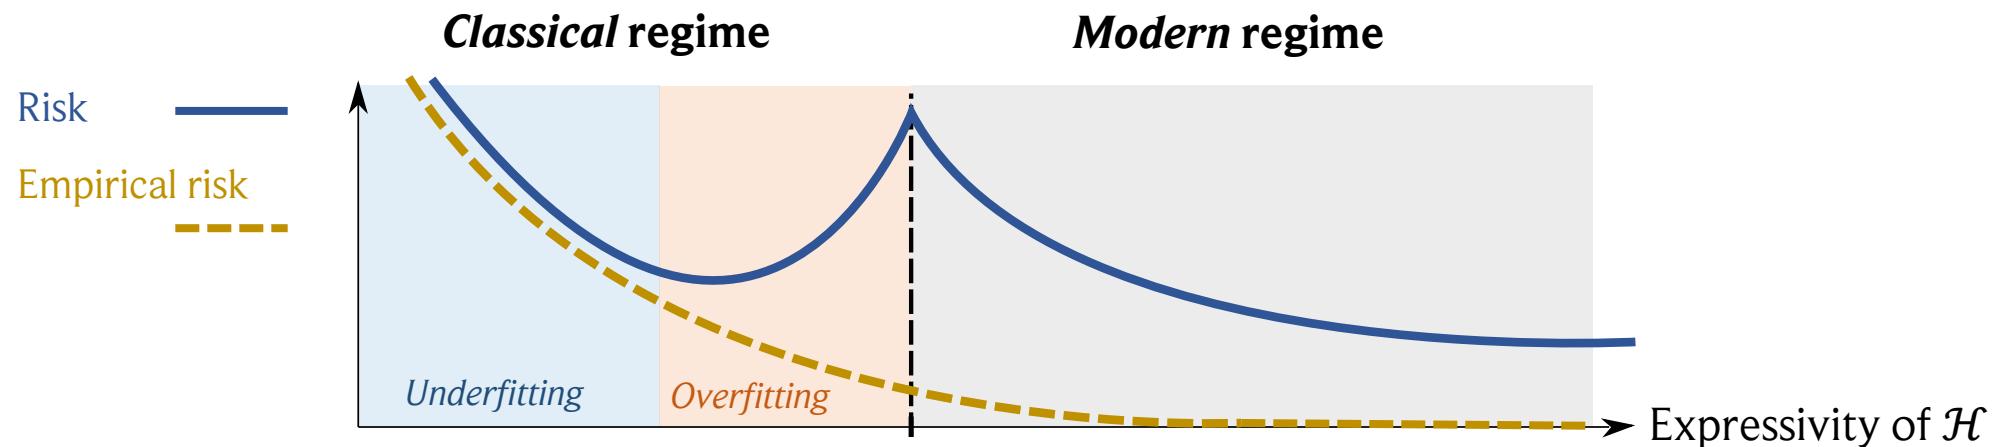
# Outro: data science in the overparametrized regime

**Classical piece of wisdom:** the more parameters your model has, the more expressive it is. However, with many more parameters than training data, it will start overfitting!



# Outro: data science in the overparametrized regime

**Classical piece of wisdom:** the more parameters your model has, the more expressive it is. However, with many more parameters than training data, it will start overfitting!

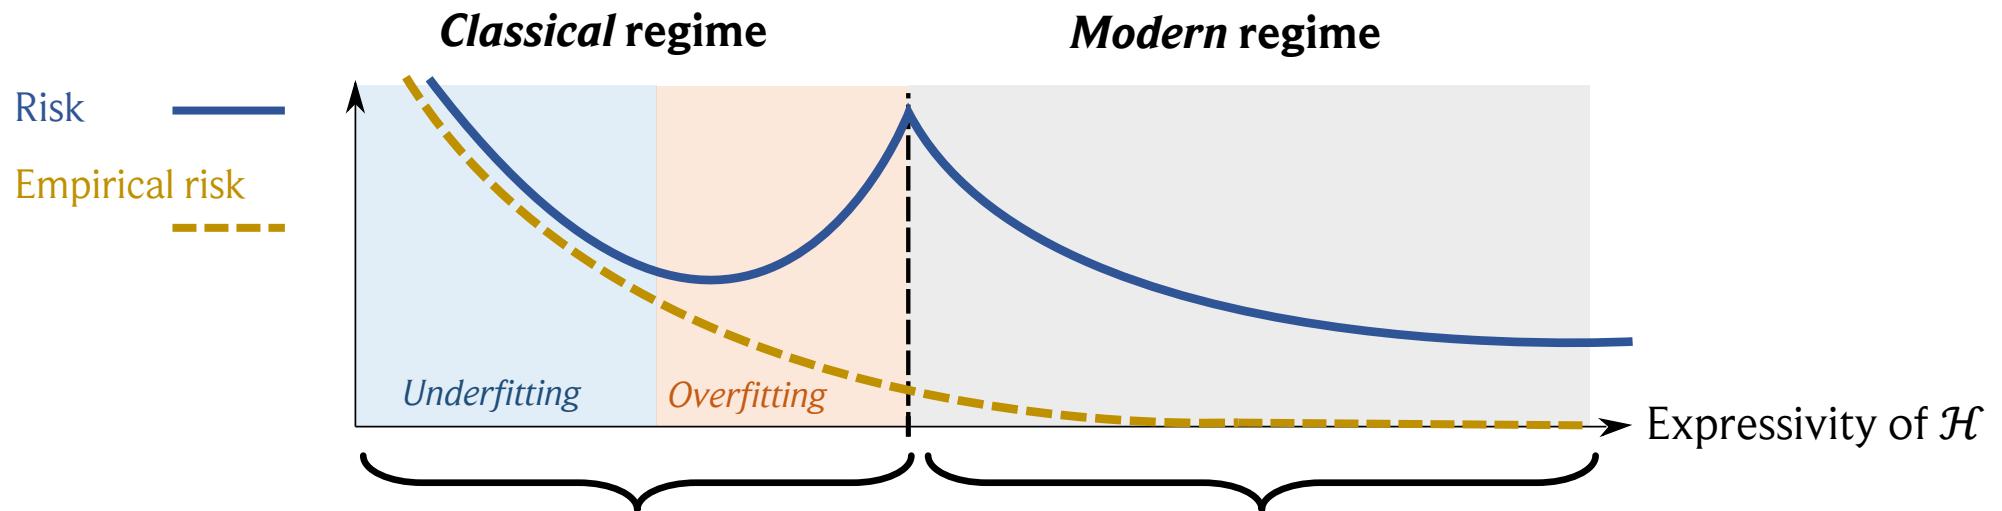


**However**, in practice we observe that highly overparametrized models can generalize well! A famous phenomena is the above curve, called “**double descent**” .

Note that this phenomenon is **not exclusive to neural networks**: one can recreate it for linear models, polynomials, regression with basis functions, etc.!

There are several explanations for this behaviour ... we will briefly touch on one of them.

# One explanation: the power of many, small parameters



Only few parameters, which potentially have to be large to fit the data. We have:

$$\mathcal{G} \leq (n/\kappa)^n \cdot (2 C_\phi B d)^{n \cdot L}$$

**Generalization bound depends on # parameters!**

Many parameters, so they can be quite small\*  
(large values are obtained by summing up many inputs!)

Thus, we can choose  $B \leq (d C_\phi)^{-1}$ , leading to  $C \approx 1$ .  
The hypothesis set of  $C$ -Lipschitz functions (here:  $C = 1$ ) has a covering number\*\* that only depends on the # input features  $d_0$ :

$$\log \mathcal{G} \leq \alpha \kappa^{-d_0} \quad \text{with constant } \alpha > 0$$

**Generalization bound is constant w.r.t. # parameters!**  
We see a decrease since the empirical risk still improves.

\* Note that most types of regularization ensure small weights ( $L1, L2$ , etc.)! We also often initialize neural networks with weight scales that depend on the width!

\*\* Both bounds shown here are valid for both regimes. We get the “double” curve by choosing the bound that is smaller!